

Udacity's Deep Reinforcement Learning Nanodegree Project 2 Report:

December 24, 2018

1 Introduction

2 Learning Algorithm

The Deep Deterministic Policy Gradient (DDPG) algorithm [1, 2] is implemented in this project:

- 1: Initialize replay memory D with capacity N
- 2: Initialize critic network \hat{q} with random weights w^q
- 3: Initialize actor network μ with weights w^μ
- 4: Initialize target critic weights $w^{q-} \leftarrow w^q$
- 5: Initialize target actor weights $w^{\mu-} \leftarrow w^\mu$
- 6: **for** the episode $e \leftarrow 1$ to M **do**
- 7: Initialize a random process \mathcal{N} for action exploration
- 8: Receive initial input state S
- 9: **for** time step $t \leftarrow 1$ to T **do**
- 10: Choose action $A = \mu(S|w^\mu) + \mathcal{N}$
- 11: Take action A , observe reward R , and next input frame S'
- 12: Store experience (S, A, R, S') in replay memory D
- 13: $S \leftarrow S'$
- 14: Obtain minibatch of tuples (s_j, a_j, r_j, s_{j+1}) from D of size K .
- 15: Set target $y_j = r_j + \gamma \hat{q}(s_{j+1}, \mu(S', w^{\mu-}), w^{q-})$
- 16: Update w^q : $\Delta w^q = -\alpha \frac{1}{N} \sum_j (y_j - \hat{q}(s_j, a_j, w^q)) \nabla_{w^q} \hat{q}(s_j, a_j, w^q)$
- 17: Update w^μ with policy gradient:

$$\nabla_{w^\mu} J \approx \frac{1}{N} \sum_i \nabla_a \hat{q}(s_i, a, \mu(s_i)) \nabla_{w^\mu} \mu(s_i | w^\mu)$$

- 18: Soft update w^{q-} : $w^{q-} \leftarrow (1 - \tau)w^{q-} + \tau w^q$
- 19: Soft update $w^{\mu-}$: $w^{\mu-} \leftarrow (1 - \tau)w^{\mu-} + \tau w^\mu$

3 Results and Plots of Rewards

4 Ideas for Future Work

References

- [1] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, sep 2015.
- [2] Udacity. Deep Reinforcement Learning Nanodegree Course Material, 2018.